

01100101
01111000
01110000
01101100
01101111
01110010
01100101
00001010



Agentic AI

Eight Insights for Business Leaders

Eight questions leaders should be asking themselves and their organisations about the architecture, evaluation, governance, and strategic horizon of agentic systems.

01

INSIGHT

Beware the Seduction of the Demo



THE QUESTIONS

"Are we mistaking a working demo for a working deployment?"

"What would it take to honestly stress-test the gap between the two?"

HOW LEADERS SHOULD PROCEED

- Reframe scaling: not enlarging a pilot: transitioning to an operational context no one has fully tested.
- Map the corner cases before deployment: where do edge conditions, hallucinations, and error propagation live?
- Build a deliberate demo-to-production checklist; assume the distance is greater than the prototype suggests.
- Resist resource allocation decisions made on the basis of an impressive prototype alone.

02

INSIGHT

The Workflow: Not the Task: is the Unit of Value



THE QUESTIONS

"Where in our organisation is the value lost in the connective tissue between tasks rather than in the tasks themselves?"

"Are we automating individual steps when we should be redesigning the choreography?"


HOW LEADERS SHOULD PROCEED

- Shift the question from 'how do we add AI to our existing process?' to 'if we were designing this process today, what would it look like?'
- Identify one high-value workflow currently held together by manual handoffs, sequencing, and judgment calls.
- Treat workflow redesign as a leadership decision, not a technology one — automation removes humans; redesign reimagines the choreography.
- Measure value at the workflow level (end-to-end coherence), not at the task level (a single step made faster).

03

INSIGHT

Autonomy is a Dial, Not a Switch



THE QUESTIONS

"Where on the autonomy spectrum should each of our deployments sit, and why?"

"Are we calibrating by stakes and reversibility, or by enthusiasm and convenience?"

HOW LEADERS SHOULD PROCEED

- Recognise the enterprise needs a portfolio of positions on the dial: different agents, different settings.
- Calibrate by context: customer-facing agents toward the constrained end; internal research agents further toward the autonomous end.
- Define explicit thresholds tied to value, jurisdiction, and reversibility — not blanket policies.
- Treat autonomy creep as the silent risk: each expansion should be a deliberate decision, not a drift.

04

INSIGHT

Evaluate Against the Current Process, Not Against Perfection



THE QUESTION

"What is our honest baseline — and are we judging the agent against it, or against an imagined ideal?"

HOW LEADERS SHOULD PROCEED

- Document the current process candidly: error rates, cycle time, audit-trail gaps, inconsistencies.
- Reframe the evaluation question: 'does this agent, with these safeguards, produce better outcomes than the process it replaces?'
- Value traceability and consistency, not just accuracy: **an 85% agent with full audit logs may beat an 82% human process.**
- Build measurement into deployment from day one (continuous evaluation, not a pre-launch gate).

05

INSIGHT

Architect Safety into the System Around the Model



THE QUESTIONS

"Where does our safety logic live: inside the model's reasoning, or in the surrounding architecture?"

"What survives a determined prompt injection?"

HOW LEADERS SHOULD PROCEED

- Move critical constraints out of prompts and into application logic: the model cannot talk its way past hard-coded rules.
- Layer the defences: hard-coded constraints, output verification, reversibility by design, escalation architecture.
- Apply the irreversibility principle: never make an irreversible decision faster than your ability to detect it was wrong.
- Design escalation paths that route by confidence and authority, verified by code rather than by language.

06

INSIGHT

The Binding Constraint is Organisational, Not Technical



THE QUESTIONS

"Are we treating this as a technology programme when the harder work is identity, incentives, and culture?"

"What is our honest position on the AI-first to AI-native gradient?"

HOW LEADERS SHOULD PROCEED

- Acknowledge that scaling agentic AI touches identity: **your most capable people may become your most significant source of resistance.**
- Pursue a dual strategy: top-down sponsorship for direction; bottom-up experimentation for conviction.
- Distinguish AI-first (it appears in strategy documents) from AI-native (people think and act with AI as default).
- Invest in change management, role design, and 'agent manager' capabilities — not just tools and licences.

07

INSIGHT

Compressed Cycles are Discipline, Not Recklessness



THE QUESTIONS

"Is our six-month evaluation cycle itself a source of risk?"

"Is our tempo matched to the rate at which the underlying technology is changing?"

HOW LEADERS SHOULD PROCEED

- Treat lengthy evaluation periods as a form of risk: by the time you finish, **the thing you evaluated is no longer the thing in front of you.**
- Adopt a three-week MVP discipline: fishfood with five people, then learn before scaling.
- Hold scope tight rather than slowing down — you can move fast when your blast radius is small.
- Build cultural agility around tech stacks; don't marry a vendor in a market where features converge in months.

08

INSIGHT

Methodology is Becoming the New Source Code



THE QUESTIONS

"What in our organisation is methodology rather than software — and could it be encoded as an agentic workflow?"

"Are we investing in assets that appreciate or depreciate in an AI-native world?"

HOW LEADERS SHOULD PROCEED

- Audit your methodologies: playbooks, frameworks, decision rules, as candidate intellectual assets to encode.
- Recognise the strategic shift: companies whose core asset is a methodology are quietly becoming software companies.
- Invest in appreciating assets: structured knowledge, clean data, documented processes, system ontologies.
- Watch for the depreciating ones: rigid custom code, tightly coupled vendor dependencies, manual processes that resist redesign.

THE THROUGH-LINE

One sentence to anchor the session

"The difference between a language model that generates text and an agent that acts in the world is not a matter of scale, but of scaffolding."

Insights 1–3

What agentic AI actually is.

Insights 6–7

What changing the organisation requires.

Insights 4–5

How to evaluate and govern it.

Insight 8

Where the strategic horizon is heading.

01100101
01111000
01110000
01101100
01101111
01110010
01100101
00001010

