



MIT SLOAN SCHOOL OF MANAGEMENT
MIT SCHWARZMAN COLLEGE OF COMPUTING



IMPLEMENTING AGENTIC AI **Your organizational playbook**

Table of contents

How to use this playbook	3
Module 1: The vocabulary week	4
1.1 Define agentic AI in your own words	4
1.2 Position your organization on the continuum	5
1.3 Apply the three adoption realities	6
1.4 Identify your candidate workflow	7
1.5 Running recommendation: The WHAT	8
Module 1 rubric	9
Module 2: The building week	11
2.1 Design the agent: Five building blocks	11
2.2 Decide to build, buy, or stack	12
2.3 Set the threshold map: Human-in-the-loop	14
2.4 Draft a three-week MVP plan	15
2.5 Determine who manages this agent	16
2.6 Define AgentOps: The measurement discipline	17
2.7 Running recommendation: The HOW	17
Module 2 rubric	18
Module 3: The scaling week and board recommendation	20
3.1 Position agents on the autonomy dial	20
3.2 Determine the four safeguard layers	21
3.3 Make the comparison that matters	22
3.4 Layer lines of defense at portfolio scale	23
3.5 Manage non-human identity	24
3.6 Set differentiating governance	24
3.7 Running recommendation: The SCALE	25
Board Recommendation	26
A. Body	26
B1. Value creation across five dimensions	26
B2. ROI narrative: Near term and long term	28
C. The closing paragraph: The ask and the commitment	28
D. Executive summary	29
Module 3 rubric	30

How to use this playbook

This document is the critical learning asset of the program. You will complete it in three sittings, one per module. The choices you make in Module 1 will frame the design work you do in Module 2, which in turn will frame the scaling and governance choices you make in Module 3. By the end of the program, this document will contain a coherent, defensible recommendation for your organization, written in your voice, for your board.

Each section of this playbook corresponds to one of the three governing questions of the program:

- **Module 1:** What is agentic AI?
- **Module 2:** How will you implement it?
- **Module 3:** How will you scale it responsibly?

The “running recommendation” paragraphs in each are the through-line of the playbook. At the end of Module 3 you will assemble these, in order, as the body of your final board recommendation. Throughout, aim for the tone of an executive note: considered, candid, and written for a colleague who is smart but unfamiliar with your specific context.

Some questions provide an optional scaffold. You are not required to use this scaffold. It is offered as a starting point, in the spirit of writing a first draft you can revise. Strong paragraphs often depart from it.

Take a moment to fill in the details below. You will refer back to this Playbook throughout the program; making it yours from the outset matters.

Name	
Organization	
Role	
Date	

Module 1: The vocabulary week

In this section, you make the foundational choice every agentic AI deployment depends on: deciding what you mean by the word, and where on the AI continuum your organization's first serious opportunity lives.

"Before we can tell the board what to do, we have to tell each other what we're talking about."

The fictional CEO of Boulton Industries

Module 1 provides three distinct registers: McDonagh-Smith's framing of agentic AI on the AI continuum, Jacob Andreas's technical introduction to the building blocks of an agent, and Kanchana Patlolla's practitioner perspective on what enterprise adoption looks like. Now, turn that vocabulary into a position. Not a finished plan – a position. The work is to decide, in language you can defend to a colleague, what agentic AI is, where it sits on your organization's horizon, and which workflow you intend to make the center of your learning.

1.1 Define agentic AI in your own words

"An agent is something my people are going to have to learn to manage. Not use. Manage. That's the definition I'm holding in my head."

Priya Kapoor, fictional Chief People Officer, Boulton Industries

1. In one sentence, write your own working definition of agentic AI. Aim for the sentence you would say at the start of an executive meeting in your organization.

Add your one-sentence definition here:

2. Now reflect: What does your definition prioritize: the technical architecture, the decision authority, the unit of work, the operational queue, or the people who manage it? What does that priority reveal about how you are inclined to frame agentic AI in your organization?

Add your three to four sentences response here:

3. Where in your organization might colleagues currently be using the word “agent” in materially different ways from your definition? Name two specific contexts (a function, a project, a team) and describe the divergence.

Add your two specific contexts, in one short paragraph each, here:

1.2 Position your organization on the continuum

“The difference between being AI-first and AI-native is the difference between a strategy and a culture, and agentic AI requires the second.”

Adapted from Kanchana Patlolla, Module 1 Unit 3

1. Locate your organization on the AI continuum. Is it primarily extracting value from classical machine learning, generative AI as a productivity tool, early agentic deployments, or some combination of these? Be specific. Name the systems, processes, or functions that anchor your assessment.

Add three to five sentences, with concrete examples, here:

2. Patlolla draws a distinction between AI-first organizations (where AI appears in strategy documents and initiative lists) and AI-native organizations (where AI is the default mode of thinking and acting). Where does your organization truly sit on that spectrum? What evidence supports your assessment?

Add three to five sentences here:

3. If your organization made a serious commitment to agentic AI in the next 12 months, which of the earlier waves (classical ML readiness, generative AI fluency, data infrastructure, change capacity) is the one most likely to constrain you? Why?

Add two to three sentences here:

1.3 Apply the three adoption realities

First: Speed of iteration matters more than perfection of plan. The technology landscape is moving so fast that a six-month evaluation cycle is itself a risk.

Second: The biggest barriers are organizational, not technical. Models are mostly capable enough; organizations are mostly not ready.

Third: Existing metrics will not capture the value agentic systems create. The KPIs you have today were likely built for a different unit of work.

1. **Speed of iteration:** Where in your organization is the current evaluation or procurement cadence likely to be a strategic risk for agentic AI specifically? Name a process or governance step where the cadence needs to compress, and estimate by how much.

Add two to four sentences, with a concrete cadence implied, here:

2. **Organizational barriers:** If your first serious agentic deployment fails or underperforms, what are the two most likely organizational reasons? Be specific: silos, incentive misalignment, data ownership, change capacity, identity issues, governance immaturity and so on.

Name two reasons, writing a short paragraph for each, here:

3. **New metrics:** Pick one workflow in your organization where, if an agent did the work, the existing metric (cycle time, cost saved, headcount avoided) would understate the value created. What new metric is missing?

Add three to five sentences, naming the workflow and the missing metric, here:

1.4 Identify your candidate workflow

“Don’t bring the safest one. Don’t bring the most ambitious one. Bring the one that will teach us the most.”

The fictional CEO of Boulton Industries

A strong candidate workflow tends to share several characteristics. It is multi-step rather than a single prompt. It crosses functions, or at least systems. It involves judgement at one or more steps, not pure mechanical processing. It has a measurable baseline today (cycle time, error rate, cost per transaction). And, importantly, it surfaces the kind of decisions you most need to learn to make as an organization, about thresholds, safeguards, and the new role of an agent manager. The point is not to pick the workflow most likely to succeed. The point is to pick the one that will teach you the most about how to make agentic decisions well.

1. Name your candidate workflow in a single phrase, as you would on a slide. Then describe it in two or three sentences for someone outside your organization.

Add one name and two to three sentences of description here:

2. Where in the workflow does the human judgement currently live? Name the specific decision points and what makes each one a judgement rather than a calculation.

Add two or three decision points, in one or two sentences each, here:

3. What is the measurable baseline today (cycle time, error rate, cost per transaction, throughput per quarter)? What numbers exist that you could compare an agent’s performance against?

Add two to four numerical or quasi-numerical baselines here:

4. What will this workflow teach your organization that an easier or safer choice would not? Be specific about the kinds of decisions it will force you to make for the first time.

Add three to five sentences here:

1.5 Running recommendation: The WHAT

This paragraph captures the WHAT of your recommendation. It does three things:

1. Names the strategic posture you are proposing your organization take toward agentic AI;
2. Identifies the specific workflow you intend to make the center of your initial learning;
3. Grounds both in honest reference to where your organization actually is on the AI continuum and the three adoption realities. It is the paragraph your board reads before any of the technical or governance detail.

Optional scaffold:

[Your organization] should respond to agentic AI by [a strategic posture, such as “deliberately and at measured pace”; “aggressively in two specific functions”; “through a portfolio of focused experiments” etc.], focusing initial learning on [your candidate workflow], because [your honest assessment of where the organization is, and what the three adoption realities imply for it].

1. Write a coherent WHAT paragraph. Do not use bullet point. This text will sit at the top of a board document.

Add four to seven sentences here:

2. Read your paragraph back. If a member of your board read only this paragraph and nothing else, what would they think you are recommending? Is that what you want them to conclude? Revise accordingly.

Add a short revision note to yourself, if any, here:

Module 1 rubric

	Criteria not met	Criteria met	Good	Exceptional
Sections 1.1–1.3 <i>Vocabulary and context</i>	No submission. OR The response fails to define agentic AI or to place it on the AI continuum.	The response attempts to define agentic AI and to place it on the AI continuum, but insight into the three adoption realities are vague.	The response provides sound insight, defining agentic AI, placing it on the AI continuum, and applying the three adoption realities.	The response provides thorough insight, defining agentic AI clearly, placing it on the AI continuum, and rigorously applying the three adoption realities. The response demonstrates a high level of expertise in the module's content.
Section 1.4 <i>Candidate workflow</i>	No submission. OR The response does not demonstrate insight into the candidate workflow.	The response attempts to provide insight into the candidate workflow.	The response provides sound insight into the candidate workflow.	The response provides thorough insight into the candidate workflow and demonstrates a high level of expertise in the module's content.
Section 1.5 <i>Running recommendation: WHAT</i>	No submission. OR The response does not provide a sound running recommendation.	The response attempts to provide a sound running recommendation.	The response provides a sound running recommendation.	The response provides an insightful running recommendation and demonstrates a high level of expertise in the module's content.

	Criteria not met	Criteria met	Good	Exceptional
<i>Clarity and logic of the submission's structure</i>	No submission. OR The response is unstructured, unclear, and lacks succinctness.	The response is logically structured, clear enough to comprehend, and mostly succinct.	The response is well structured in terms of logic, clarity, and succinctness.	The response is exceptionally well structured in terms of logic, clarity, and succinctness.

Module 2: The building week

In this section, you turn the candidate workflow you identified into a designed agent, making the architectural, adoption, and operational decisions that determine whether what you build will work in the world.

“An agent is a language model plus scaffolding.”

Adapted from Jacob Andreas, Module 2 Unit 1

2.1 Design the agent: Five building blocks

Andreas argues that building a language model agent requires answering five questions, and that every design decision downstream is a commitment on one of these five axes. The discipline of working through them in sequence is what separates a real agent from a wished-for one.

“What turns a language model into an agent is not a new kind of model but a change in interpretation: The same text output, interpreted as a program rather than as a reply, becomes an action.”

Adapted from Jacob Andreas, Module 2 Unit 1

Complete each row of Table 1 for your candidate workflow. Be concrete. Where you do not know the answer yet, name the open question – that is itself a useful output.

Table 1: Five building blocks.

Building block	The question to answer	Your design decision
Perception	How will the agent observe the world? What inputs, in what formats, from what sources?	
Action	What is the agent allowed to do? Which tools, APIs, or systems can it invoke, and with what side effects?	

Planning	How will the agent reason about multi-step tasks? Single-step, chain-of-thought, decomposition into sub-agents?	
Memory	What does the agent need to remember between turns, between sessions, between deployments?	
Safety	How is risk architected into the surrounding system, not the model? What does the system do that the model cannot be trusted to do?	

1. Andreas posits that you should not rely on the language model itself to manage risk. Instead, you architect the safety into the surrounding system. Looking at your Safety row above, where in your design is the riskiest place you are currently relying on the model itself to make a judgement that it shouldn't?

Add one specific place and one or two sentences on why moving it outside the model would matter:

2.2 Decide to build, buy, or stack

Patlolla addresses the strategic posture question every leader will face: At each layer of the agentic stack (model, framework, platform), do you build, buy, or stack a combination? Her answer is consequential and disciplined. The decision is driven by the autonomy you need and the skill set you have, not by technology preference.

"It doesn't have to be one or the other. It could be a combination of all three in your enterprise."

Kanchana Patlolla, Module 2 Unit 2

For your candidate workflow, complete the build/buy/stack matrix in Table 2. In each cell, note the choice you are inclined toward, and one or two sentences justifying it for this specific deployment.

Table 2: Build, buy, stack matrix.

Layer	Build	Buy	Stack (combination)
Model			
Framework			
Platform			

1. Patlolla's underlying argument is that the choice at each layer is a function of two variables: the autonomy you need, and your internal skill set. For your candidate workflow, which layer is most likely to expose a skill-set gap in your organization? What would closing that gap look like (building internal capability, hiring, partnering, etc.)?

Add three to four sentences here:

2.3 Set the threshold map: Human-in-the-loop

"I don't care about a \$5 fraud, let it go. That's something we've always done, even pre-generative AI: We've created thresholds and handled those cases."

Kanchana Patlolla, Module 2 Unit 2

For your candidate workflow, draw the threshold map. In Table 3, identify three decision types the agent will encounter. For each, name the threshold at which the agent acts autonomously, the threshold at which it escalates, and the threshold above which a human must originate the decision.

Table 3: Threshold map.

Decision type	Agent acts autonomously when...	Agent escalates to human when...	Human originates the decision when...
1.			
2.			
3.			

- Looking at your threshold map, which threshold are you least confident about (where you genuinely do not yet know where the line should be)? What experiment would help you find out?

Add one specific threshold and one or two sentences on the experiment that would calibrate it here:

2.4 Draft a three-week MVP plan

Patlolla's recommended cadence is three weeks (and five people) from design board to first internal release. The discipline is not in the speed but the scope. A three-week MVP is honest only when the scope is held tight. Pick one slice of your candidate workflow that can credibly be agentified in three weeks.

"Three weeks is even a lot, because you can build a prototype for a complex vision in less than a weekend."

Kanchana Patlolla, Module 2 Unit 2

1. **MVP scope:** What is the smallest, end-to-end slice of your candidate workflow that an agent could plausibly handle? Describe it the way you would describe it to the engineer who would build it. Be specific about start and end points.

Add two to four sentences here:

2. **MVP audience:** Who are the five (or fewer) people who would use the MVP first? Name the role and describe the type of feedback you would solicit from each.

Add up to five named roles, one line each, here:

3. **What to leave out:** The hardest part of a three-week MVP is what you leave out. Name two specific aspects of your candidate workflow that you are deliberately deferring beyond the MVP, and the trigger condition that would bring each one back into scope.

Add two deferred aspects and triggers, in one or two sentences for each, here:

2.5 Determine who manages this agent

"Every agent needs a manager."

Adapted from Kanchana Patlolla, Module 2 Unit 2

1. Who, by name or by role, will be the day-to-day manager of your agent? Why this person? What do they bring that the engineering team would not?

Add three to five sentences here:

2. What does the agent-manager role need to do that this person's current role does not? Name three specific responsibilities that are new (not adapted versions of existing ones).

Add three specific new responsibilities, in one short paragraph each, here:

3. Will this be an add-on to an existing role, or a new role? If an add-on: How will their existing workload be adjusted? If new: What is the path to formalizing this position in your organization?

Add two to four sentences here:

2.6 Define AgentOps: The measurement discipline

Patlolla's second high-conviction framing is the move from CI/CD/CT/CM (continuous integration, delivery, training, monitoring – the MLOps inheritance) to what the field is calling AgentOps. The shift is to continuous evaluation as a first-class deployment concern, with skew detection built into the architecture.

"The dashboard is not a report. It is a conversation we are now having continuously with our own system."

Rebecca Holloway, fictional CFO, Boulton Industries

1. **Continuous evaluation:** What does the agent need to be evaluated against, in production, on a continuous basis? Name two metrics that should be tracked in real time, and one quality measure that requires periodic human review.

Add two real-time metrics and one quality measure with a proposed review cadence here:

2. **Skew detection:** Where in your candidate workflow is the agent's behavior most likely to drift over time? Is it because the input distribution shifts, the underlying model is upgraded, or a tool's API changes? What would early warning of that drift look like?

Add two to four sentences here:

3. **Ownership:** Who reviews the AgentOps dashboard, on what cadence, and with what authority? This is a governance question as much as an engineering one.

Add three to five sentences here:

2.7 Running recommendation: The HOW

The HOW paragraph integrates the work of this section into a single coherent statement. It tells the reader what kind of agent you are building (the building-block specifications, in compressed form), what your build/buy/stack posture is, where the human-in-the-loop thresholds sit, who will manage the agent day-to-day, and what AgentOps discipline will hold it to standard. Five things in one paragraph.

Optional scaffold:

[Your organization] will build this capability by [your build/buy/stack posture], deploying an agent with [a one-line summary of the building-block decisions] and human-in-the-loop thresholds at [the consequential decision points], managed day-to-day by [the agent-manager role], and held to standard through [your AgentOps cadence].

1. Write your HOW paragraph (again, do not use bullet points).

Add five to eight sentences here:

2. Read your two running recommendation paragraphs together. Do they cohere? If a colleague read both as a single argument, would they understand how the design choices in the second paragraph follow from the strategic posture in the first? Note any revision your earlier paragraph now needs in light of the work you have just done.

Add a short revision note to yourself, if any, here:

Module 2 rubric

	Criteria not met	Criteria met	Good	Exceptional
Sections 2.1, 2.2, & 2.4 <i>Design and architecture</i>	No submission. OR The response fails to complete the five building blocks table, the build, buy, or stack matrix, or the three-week MVP plan.	The response attempts to complete the five building blocks, the build, buy, or stack matrix, and the three-week MVP plan.	The response provides sound insight on the five building blocks, a clear build, buy, or stack approach, and a realistic three-week MVP plan.	The response provides thorough insight on the five building blocks, a clear build, buy, or stack strategy, and a specific three-week MVP plan. The response demonstrates a high level of expertise in the module's content.
Sections 2.3, 2.5, & 2.6 <i>Operations and management</i>	No submission. OR The response does not demonstrate insight into the thresholds and management plans.	The response attempts to provide insight into the thresholds and management plans.	The response provides sound insight into the thresholds and management plans.	The response provides thorough insight into the thresholds and management plans and demonstrates a high level of expertise in the module's content.

	Criteria not met	Criteria met	Good	Exceptional
Section 2.7 <i>Running recommendation: HOW</i>	No submission. OR The response does not provide a sound running recommendation.	The response attempts to provide a sound running recommendation.	The response provides a sound running recommendation.	The response provides an insightful running recommendation and demonstrates a high level of expertise in the module's content.
<i>Clarity and logic of the submission's structure</i>	No submission. OR The response is unstructured, unclear, and lacks succinctness.	The response is logically structured, clear enough to comprehend, and mostly succinct.	The response is well structured in terms of logic, clarity, and succinctness.	The response is exceptionally well structured in terms of logic, clarity, and succinctness.

Module 3: The scaling week and board recommendation

“Scaling is not enlarging a pilot. It is transitioning from a controlled environment to an operational context that no one has fully tested.”

Paul McDonagh-Smith, Module 3 Unit 1

Module 3 shifts the level of analysis from the agent to the system. In this section, you complete the work of the program: positioning your agent on the autonomy dial, designing the safeguard architecture that makes it durable, and articulating the governance posture that will let your organization deploy not one agent but many –responsibly, measurably, and over time.

3.1 Position agents on the autonomy dial

“Autonomy is not a switch you flip. It is a dial you calibrate.”

Paul McDonagh-Smith, Module 3 Unit 1 Video 2

The critical design decision is where you position each agent, for each task, in each context:

- **Constrained:** Agent only returns pre-verified responses.
 - **Bounded:** Agent operates within a narrow, validated decision space with near-deterministic rules.
 - **Supervised:** Agent acts but every consequential action is reviewed before execution.
 - **Delegated:** Agent acts autonomously within thresholds; humans review by exception.
 - **Autonomous:** Agent acts and self-monitors; humans review only on flagged anomalies.
1. For your candidate workflow, determine autonomy for each of the three core decision types in your Module 2 threshold map and explain your placement in one or two sentences each.

Add three placements with reasoning here:

2. Where on the dial does the agent currently sit by default in your design? Where do you imagine it sitting in 12 months, if it performs as intended? What conditions would have to be met for it to migrate from the first position to the second?

Add three to five sentences here:

3.2 Determine the four safeguard layers

For your candidate workflow, complete the four-layer safeguard in Table 4. Each row asks you to translate the architectural principle into a safeguard specific to your design. Where a layer does not yet exist in your design, name what it would take to add it.

Table 4: Four-layer safeguard.

Safeguard layer	Architectural principle	Concrete safeguard in your design
Hard-coded constraints	Non-negotiable boundaries built into the system architecture itself, executing independently of the model's reasoning. The model cannot talk its way past them.	
Output verification	Before an agent's output reaches a user or triggers a downstream action, it is checked against formal rules, consistency criteria, or a second model acting as reviewer. The two-person rule, in software.	

Reversibility by design	Wherever possible, agent actions are designed to be reversible. Reversibility limits the blast radius of an error; it does not prevent the error itself.	
Escalation architecture	Every agent has a clear, tested escalation path. This is a mechanism for recognizing the boundary of its confidence or its authority and routing the decision to a human.	

1. Look across your four layers. Which is the strongest in your current design? Which is the weakest? What single intervention would most improve the weakest layer?

Add three or four sentences here:

3.3 Make the comparison that matters

When evaluating the risk of an agentic system, the comparison is not agent versus perfection, it is agent-with-safeguards versus the current process. The question to put to your board is not “Does the agent make errors?” but “Does this agent, with these safeguards, produce better outcomes than the process it replaces, on the dimensions that matter?”

“We are replacing processes which normally humans do, and they also make mistakes.”

Tim Kraska, Module 3 Unit 2

1. Sketch the head-to-head comparison for your candidate workflow. On the one hand, the current human-driven process (accuracy, cycle time, consistency, audit trail, error patterns, cost). On the other, the agent-with-safeguards. Where the projection is uncertain, say so.

Add four to six sentences, or a structured comparison, here:

2. On which dimension is the agent likely to be worse than the current process, at least at first? How will you defend that trade-off to your board, and what is the path to closing the gap?

Add two to three sentences here:

3.4 Layer lines of defense at portfolio scale

Most organizations that scale agentic AI successfully will not deploy one agent; they will deploy ten, or twenty, or a hundred, each calibrated differently. At that scale, governance cannot be designed agent-by-agent. It has to be infrastructural.

1. **Foundation layer: Data and model governance.** What policies, processes, or infrastructure does your organization already have in place that would extend to a portfolio of agents, and where are the most significant gaps? Be specific.

Add three to four sentences here:

2. **Operational layer: Runtime guardrails.** What input/output moderation, prompt routing, and use-case-specific constraints would need to be standardized across your agent portfolio? Who would own them?

Add four to six sentences here:

3. **Transparency layer: The AI bill of materials.** For your candidate workflow, sketch the components that would appear on its AI bill of materials (training data provenance, model architecture, version history, third-party components, tool dependencies). Where is the documentation thin?

Add five to seven sentences, or a structured list, here:

3.5 Manage non-human identity

1. For your candidate workflow, define the agent's identity. What is the agent allowed to access? What is it explicitly not allowed to access? Which tools, APIs, and systems does it have permissions for? What is the principle of least privilege, applied here?

Add five to seven sentences here:

2. How will the agent's permissions be reviewed, by whom, and on what cadence? Under what conditions would you revoke or restrict permissions? Treat this as you would when considering a new senior employee.

Add three to five sentences here:

3.6 Set differentiating governance

1. For your organization, where is the strongest case that governance investment in agentic AI would create competitive advantage rather than simply protective hygiene? Name the specific market, customer segment, or regulatory environment where this would land most clearly.

Add four to six sentences here:

2. What is the credible governance posture your organization could announce externally in the next 12 months? Describe the truth behind that announcement, not the aspirational marketing.

Add three to five sentences here:

3.7 Running recommendation: The SCALE

This paragraph integrates the work of this section into a single coherent statement: the autonomy posture, the safeguard architecture, the layered lines of defense at portfolio scale, the identity discipline, and the governance-as-differentiator framing. It closes with an explicit ask: telling your board not only what you intend to do, but how you intend to do it responsibly, and what you need from them to do it well.

Optional scaffold:

[Your organization] will scale this responsibly by [autonomy posture and safeguard architecture], governed through [the layered lines of defense and identity disciplines], and treating governance as [the differentiator framing specific to your market or regulatory context]. To do this we are asking the board for [the specific investment, mandate, sponsorship, governance authority], in return for [the specific commitments your team will make].

1. Write your SCALE paragraph.

Add five to eight sentences here:

2. Now read all three running recommendation paragraphs together. Does the argument hold? Is the chain of reasoning visible? Where does it strain? Make any final revisions to the earlier paragraphs you now think they need.

Add a short revision note, plus any edits you make above, here:

Board Recommendation

This assembles your work throughout this program into an artefact you can take to your board. Though it may seem counterintuitive, build the recommendation in the order presented:

A. Body

Compile your three running recommendation paragraphs, making whatever adjustments are needed for a coherent argument.

Add your running recommendations here:

B1. Value creation across five dimensions

Complete each row of Table 5 with the specific (and, where possible, measurable) value your candidate workflow's agentic version would create. Be honest where the value is small, ambiguous, or speculative – your board will respect candor and suspect overclaim.

“When an agent resolves a customer issue end-to-end, the value is not only in the minutes saved, it is in the strength or weakness of the lifetime relationship that follows.”

Adapted from the Module 1 Boulton podcast

Table 5: Five dimensions.

Dimension	What this dimension is asking	Near-term value (12 months)	Long-term value (3-5 year)
Efficiency	Reduction in time, cost, or cognitive load required to complete the work and where in the workflow that reduction is realized		

Quality	Improvement in accuracy, consistency, traceability, or completeness, including the kinds of error the agent makes less often than humans, and the kinds it may make more often		
Resilience	Capacity to handle volume spikes, staff turnover, knowledge loss, or disruption, and what becomes more recoverable when the agent is in place		
Speed	Reduction in end-to-end cycle time and the second-order effects of that reduction on customer experience, market response, or downstream processes		
Innovation	What new work becomes possible because the agent has freed capacity, opened access, or restructured the unit of effort –		

	often the largest long-term value, and the easiest to understate		
--	--	--	--

B2. ROI narrative: Near term and long term

The five-dimension table is the structure. The ROI narrative is the prose. Boards read tables for orientation, but they decide on the strength of the narrative that surrounds them.

1. Near-term ROI narrative (next 12 months). What will the agent demonstrably deliver in its first year, in measurable terms with the caveats your board will expect? Reference the value table where useful, but write in connected prose.

Add four to six sentences here:

2. Long-term ROI narrative (three to five years). What does sustained investment in this capability, and the portfolio of agents that follows from it, make possible at the operating-model level? Be concrete: name the new units of work, the new market positions, and the new organizational capabilities.

Add four to six sentences here:

C. The closing paragraph: The ask and the commitment

Every board recommendation needs a closing paragraph that brings the argument home. Remember that vague asks produce vague responses; specific asks produce decisions.

1. Write the closing paragraph of your recommendation. Name the ask (mandate, investment, sponsorship, authority); name the commitment (cadence, governance discipline, transparency, exit conditions). Close with a single sentence that names what the organization gains by acting on this now rather than later.

Add your closing paragraph in three to five sentences here:

D. Executive summary

Written last but placed first. The Executive Summary is the single page your board reads before they read anything else, and decides whether they read the rest. It is the compressed statement of what you are recommending, why, and on what terms.

“Boulton Industries should deploy agentic AI, deliberately and at measured pace, in a manner that treats governance, safeguard architecture, and organizational readiness as the central design variables rather than downstream constraints.”

The fictional CEO, Boulton Industries

1. Write the single opening sentence of your Executive Summary. The sentence should state, in one breath, the recommendation itself – the strategic posture, the scope, and the qualifier that signals seriousness.

Add one sentence here:

2. Compress your WHAT paragraph into three to four sentences for the Summary. Strip detail; keep argument.

Add three to four sentences here:

3. Compress your HOW paragraph into three to four sentences. Name the design decisions in compressed form; do not list them.

Add three to four sentences here:

4. Compress your SCALE paragraph into three to four sentences. Include the explicit ask, in one phrase.

Add three to four sentences here:

5. Write one closing sentence. Express what the organization gains by acting now in the most direct language you can find.

Add one sentence here:

You have now written your Executive Summary. By convention, it sits at the front of the final document, between the title page and the body. When you assemble the final version of the recommendation outside this playbook, place this Summary first, followed by the body, then the value creation table and ROI narrative, then the closing paragraph.

Module 3 rubric

	Criteria not met	Criteria met	Good	Exceptional
Sections 3.1–3.3 <i>Safeguards and autonomy</i>	No submission. OR The response fails to apply the autonomy dial and safeguard layers.	The response attempts to apply the autonomy dial and safeguard layers.	The response provides sound placement on the autonomy dial, and concrete applications of the safeguard layers.	The response provides thorough, nuanced insight into calibrating autonomy and layered safeguards. The response demonstrates a high level of expertise in the module's content.
Sections 3.4–3.6 <i>Portfolio governance</i>	No submission. OR The response does not demonstrate insight into the lines of defense or identity management.	The response attempts to provide insight into the lines of defense and identity management.	The response provides sound insight into the lines of defense and identity management.	The response provides thorough insight into the lines of defense and identity management. The response demonstrates a high level of expertise in the module's content.

	Criteria not met	Criteria met	Good	Exceptional
Section 3.7 <i>Running recommendation: SCALE</i>	No submission. OR The response does not demonstrate insight into the running recommendation.	The response attempts to provide insight into the running recommendation.	The response provides sound insight into the running recommendation.	The response provides thorough insight into the running recommendation. The response demonstrates a high level of expertise in the module's content.
Board recommendation <i>Value creation and ROI</i>	No submission. OR The response does not demonstrate insight into the value creation and ROI.	The response attempts to provide insight into the value creation and ROI.	The response provides sound insight into the value creation and ROI.	The response provides thorough insight into the value creation and ROI. The response demonstrates a high level of expertise in the module's content.
Board recommendation <i>Closing commitment</i>	No submission. OR The closing lacks a specific ask.	The closing is vague and does not name specific board asks.	The closing provides a specific ask.	The closing provides a highly specific ask to drive the board towards immediate action. The response demonstrates a high level of expertise in the module's content.

	Criteria not met	Criteria met	Good	Exceptional
Board recommendation <i>Executive Summary</i>	No submission. OR The response does not provide a concise and coherent Executive Summary.	The response attempts to provide a concise and coherent Executive Summary.	The response provides a concise and coherent Executive Summary.	The response provides a motivating, concise and coherent Executive Summary and demonstrates a high level of expertise in the module's content.
<i>Clarity and logic of the submission's structure</i>	No submission. OR The response is unstructured, unclear, and lacks succinctness.	The response is logically structured, clear enough to comprehend, and mostly succinct.	The response is well structured in terms of logic, clarity, and succinctness.	The response is exceptionally well structured in terms of logic, clarity, and succinctness.