




MIT SLOAN SCHOOL OF MANAGEMENT
MIT SCHWARZMAN COLLEGE OF COMPUTING



MODULE 3 UNIT 3
Media Set Video 1 Transcript

Module 3 Unit 3 Media Set Video 1 Transcript

Introduction

PAUL MCDONAGH-SMITH: So, Glenn, a pleasure to have you here with us today. Would you mind taking a moment to introduce yourself and let us know who you are and what you do?

GLENN COHEN: Thanks, Paul. It's always a pleasure to be here with you. I'm a professor at Harvard Law School. I am a Professor of Health Law, Bioethics, and also the Deputy Dean there. And one of the areas I work on is the integration of AI into medicine, looking at legal and ethical issues raised by that integration.

Ethics by design

MCDONAGH-SMITH: So, Glenn, you work at the intersection of law, ethics, and medicine. You've seen medical AI mature into a highly regulated area of the AI landscape. For our business leaders joining us today in this module, what's the gift of that vantage point? What can they see from where you sit that's harder to see from inside their own sectors?

COHEN: Great question. So I think there's a few things to see. One is this evolution, as you say, in that now we're seeing so much more integration and it's happened extremely fast. The second is that while large language models and ChatGPT and the like are very familiar, actually much of the AI integrated into medicine are different models of AI. And so some of the lessons about large language models, some of the concerns, maybe don't carry over.

The third is, I think, just gauging risk and getting a sense of what the risks really are, right? Often, people come to me and they think the first thing I tell them is, "You're going to be sued for this reason or that reason." But I tell them, "Part of what lawyers do is we do risk assessment. We try to tell you what are the spaces that are riskier, what are less risky." And it's always a probabilistic sort of thing. But I actually think maybe one might overestimate the risk of medical AI. And then maybe the fourth thing is just to say this idea of ethics by design. Ethics is not something you think about at the end of a product life cycle when you're ready to spin it out and someone's like, "Well, what about this?"

Instead, the idea is building in the ethical analysis from the very beginning, including deciding whether something is a targetable market, for example, and exactly what kind of design. And also, what kind of engagement you're going to have with the kind of constituencies that might have some reservations about what you're doing. So those are all things that I see from where I sit that may be less familiar to some business people.

MCDONAGH-SMITH: So I think it's fascinating that this idea of ethics by design. If I'm interpreting it correctly, it's almost a positioning of providing the scaffolding at the beginning or close to the beginning of projects that allows the correct development and design to take place.

Agentic AI liability shift

MCDONAGH-SMITH: So, Glenn, I've always found your work on medical AI liability to be particularly interesting. In your opinion, do agentic systems complicate frameworks in a specific way? For example, an agent may take a sequence of actions with no single human decision point, may replan mid-task, or even act in parallel across workflows. For those leaders with us here today who are designing agentic deployments, what practices do you think would most help their organization image and implement for the near and the next futures?

COHEN: Great. So, let me say a couple of things about this. First, just to set the baseline, in medical AI before the agentic revolution, if you will, you know, we have essentially we're able to characterize what the risks look like by drawing analogies to the way medical malpractice law handles other technologies as they're introduced and how we allocate liability between developers, hospitals, and physicians. We're also able, happily, to be able to characterize the existing cases because there aren't that many. And we can look at the kinds of rationales the courts have offered. And often, actually, I think most of the physicians I talk to probably overestimate the existing risk of liability, at least as reflected by the current decisions. That's the first thing to say.

Now, the second. You introduce the agentic element to a larger extent. What happens? Well, it turns out the law actually has an entire branch of law called agency law, which is really the name agentic AI always makes me laugh because we've been thinking about principal agent tensions for a very long time. And indeed, there's a whole set of law as to when an agent can commit a principal. When are they authorized, for example, to sign a contract engaged in a deal? But in tort law, we also have a whole set of legal rules about when you are responsible for your agent's behavior. And some of this stuff, which is old school a little bit, has all this funny lingo. Like, they talk about whether the agent is on a frolic, like a delivery driver in their hour off, hit someone in the car as opposed to hitting somebody while they're driving the car.

Now that might seem very far from the AI case, but here's why it's useful to think in this sort of way, in that typically we think, is this a liability for the agent or the principal? When it comes to agentic AI, the idea of liability for the agent doesn't really make a whole lot of sense. And instead, there's a much bigger focus, I would say, and it's an emerging case law, so we'll see. There ought to be a much bigger focus about exactly how much leeway you've given to your agents, how carefully you have observed their behavior, whether you've done in a small bore way before you unleash them out into the world, and what your obligations are in terms of monitoring and correction as time goes, as well as how you communicate what the agent is and what's the agent doing to the people who are interacting to it.

So, I actually think that compared to non-agentic medical artificial intelligence, where we can really well characterize what the liability risks are, when it comes to agentic AI, we're much more focused on what kinds of constraints you've put on and what kind of design decisions you've put on at the forefront. And I suspect that the case law is going to emerge to kind of reflect that. And the way that feeds back is just to say, there are many tasks which would be much cheaper for you to do with an agent. Can you predict ahead of time in what percentage of those cases the agent will go awry, what the consequences will be?

Because you're likely to be left holding the bag as the person who's decided to employ an agentic AI design versus other kinds of designs you might have. And I do think the courts

are going to increasingly look at that design choice moment as a moment that tells us something about the liability of the tortfeasor.

MCDONAGH-SMITH: Definitely food for thought, as well as a foundation and a framework for how we might be thinking about these types of questions. So, thank you. Thank you, Glenn.

SPEAKER: Did you understand all the concepts covered in this video? If you'd like to review any of the content, please click on the chapter menu.