




MIT SLOAN SCHOOL OF MANAGEMENT
MIT SCHWARZMAN COLLEGE OF COMPUTING



MODULE 3 UNIT 3
Media Set Video 2 Transcript

Module 3 Unit 3 Media Set Video 2 Transcript

Informed consent and disclosure

PAUL MCDONAGH-SMITH: So Glenn, there's a question you've surfaced that I think is regularly absent from the enterprise AI conversation. When do you have to tell someone that they're actually interacting with or being evaluated by an agent. So, for the leaders who are here with us today, thinking about employees, customers, citizens, patients, where do you see disclosure obligations heading over the next period?

GLENN COHEN: So great. So, we do have some places where the law says you must disclose. Right? Including, for example, California has a chatbot law relating to minors where it says you have to constantly remind them that this is not a human being. Okay. But then you have a much greater swath where I say the rules, the legal rules at least, are not so clear. And the question I think is much more with my ethicist hat on, which is to say, "What is the right way to do it and design a product?" And typically, I start from a strong presumption that when it comes to interaction rather than evaluation or information, when an individual user is interacting directly with an artificial intelligence chatbot, there should be a very strong presumption of disclosure.

When I think both about the idea that people are, especially in medicine, in extremely vulnerable states and communicating information and to find out that they've been misled or they at least haven't been made clear that they're interacting with not their physician, not a nurse, not someone at the hospital, but an AI system, to me that's really ethically problematic. When it comes to the evaluation, I think it's a little bit trickier. And the circumstances I most work in are things like decision support, when an AI has been involved in helping a physician evaluate a result to try to determine, for example, whether to engage in a procedure for a cancer or watchful waiting, right?

Do you have an obligation to tell a patient that an AI was part of that evaluation? And I think what makes it tricky is there's all sorts of things that inform physician judgment, for example, that we never disclose. They don't talk about the seven patients you remind them of, back issues of an old medical journal they read, right? Just how annoyed and frustrated they are in the day that might drive their behavior. So, we don't think you have an obligation to disclose everything.

On the flip side, there's stuff that definitely has to be disclosed, like, if a substitute surgeon scrubs into an operation and you fail to disclose that, that's a clear legal and ethical violation. So how do we think about that on the spectrum? And what I normally tell people is my own take is that we should be thinking about things like materiality. Is this the kind of information, the evaluation of an AI, its involvement that's extremely important to this particular patient? Do I have reason to know that?

Stakes of the decision. Is this a high-stakes decision, like whether to get a serious surgery or not, where actually every piece of information is extremely valuable and people worry a lot about it? And then perhaps most importantly, how much reason does the physician have to trust the AI such that the patient, in trusting the physician, we can kind of carry that trust over? When it's an AI that has been cleared by a regulator or is the process of a very searching kind of internal evaluation, it looks quite different than something someone just downloaded on their phone and started using.

Right? So to me, we have to think about these variables, but it is contextual. And this is why I think working through it, but also engaging, in my case, patients, and some other people's cases, end users or customers about their feelings should they discover after the fact that an AI was involved and they were not disclosed. This kind of engagement is really helpful as part of the design process.

MCDONAGH-SMITH: And I think this is a very helpful positioning actually because this almost multidimensional type of approach that you're respecting the legal and regulatory kind of considerations here, whilst also respecting and recognizing that in the physical carbon world, as you mentioned, the disclosures at times become a question of discretion. And so the context becomes very important as well.

Bias as a pipeline problem, not a data problem

MCDONAGH-SMITH: So Glenn, a lot of enterprise conversations about AI bias, they seem to treat it as essentially a training data issue. For our leaders on this program deploying agentic systems across their organizations, what's demanded of them operationally? Who in the organization should own each of the key elements, and what questions should they be asking at each stage?

COHEN: I'm so glad you asked me this. I think one of the biggest misconceptions people have is, well, missing data bias. Right? The data is not suited to the problem. If I could just solve that, all the bias issues go away and it'll be great if you can solve that, don't get me wrong. But I want to kind of enrich the conversation just to see how many different sources of bias there are. So even just problem formulation, right, where you decide to intervene and whether you decide this to be the problem that you're going to solve, itself is driven by lots of assumptions about the population, about your culture, about implementation. Then again, we do have this question of missing data bias.

Other things to think about is kind of in the training data, how are things measured and how the measurements themselves might reflect real-world bias?

Operational playbook

COHEN: Okay. So what's the lesson from all of that? The lessons for me are, first, you really just have to decide whether this is a good task to be spending your time on. And there's lots of technical questions about that, but there's also just a question about, "Is this a high-value task and why am I spending the money on this?"

And really am I able, can I pilot and can I get a sense very early on whether I have the data I need to do this?" Then to think looking at the data you have and to ask yourselves, "What kind of biases may already have been encoded in the data that I'm not thinking about?" And here's sometimes some of the people who might be involved might actually be social scientists who worked on collecting the data, historians who know a little bit about where the data is coming from, but also members of the affected community who might be able to tell you, for example, that many women show up to the hospital and their heart disease or their hypertension is a little bit discounted, right?

Physicians, for example, in my space, who might be able to tell you, well, this has changed over time in medical practice, so the data is going to look different from different periods of time. As well as informatics people who project forward how big the differences would have

to be to make a difference. Then you get to the stage of actually designing the thing, and I think you want to think as much as possible to spot all the biases possible to have a diverse team, both diverse in terms of backgrounds and diverse in terms of education and approaches to try to issue spot as much as possible.

But to realize you're never going to catch them all. Right? The idea of eliminating bias is I think an impossibility in any moderately complex task. Instead, though, I think your goal is get rid of as many as possible as you can in the design process, but then have a process on the back end after implementation for actually testing in the real world to see, among other things, has there been a shift in the target population?

Has been there a shift in the behavior such that your model might have been perfectly good, but now is actually rubbish when it comes to actually doing what you want? And also thinking really seriously about the place of interaction, human–computer interaction, and the actual implementation. Again, taking lessons from medicine, which is where I come from. You can have a image classifier, like for radiology, for example, that does great in one hospital, one set of radiologists, and terrible in another because of the way workflow is organized, the way training is done, and the like.

So what kind of testing are you going to do for those localized conditions? What are you going to try to do to correct as you go? And that requires, again, a real commitment to a diverse set of expertise and also a program to actually spend time doing this kind of thing.

SPEAKER: Did you understand all the concepts covered in this video? If you'd like to review any of the content, please click on the chapter menu.