



MIT SLOAN SCHOOL OF MANAGEMENT
MIT SCHWARZMAN COLLEGE OF COMPUTING



MODULE 3 UNIT 3
Media Set Video 3 Transcript

Module 3 Unit 3 Media Set Video 3 Transcript

Agent-managed workforce

PAUL MCDONAGH-SMITH: Glenn, your perspective on clinical surveillance uses a phrase that has stayed with me. Mechanical managers. You proposed a clinician bill of rights for AI, a right of information, a right of participation, a right of privacy, and quality assurance. The flip side, the workers who are managed, evaluated, and surveilled by agents, is I think a question that your paper and your work raises.

In your opinion, do we need an equivalent bill of rights for the agent-oriented workforce? And if the leaders here with us were to draft one tomorrow morning for their organizations, what are the two, three, four rights that they'd find hardest to get right?

GLENN COHEN: Great question. So this is from work I did with Ravi Parikh and Ifeoma Ajunwa in the *New England Journal of Medicine*. And we take this idea, which is, Ifeoma has a book on this, about the quantified worker. The idea that if we look at the transformation of many blue-collar spaces, like Amazon fulfillment centers or long-haul truck drivers, we see an increasing use of AI for surveillance and management that often results in diminutions in discretion and also diminutions in status. And we say, well, now let's look ahead to white-collar workers and think about, especially in medicine, how this might change things.

So, I think agentic AI in some ways is, in terms of the labor market disruption and the effect on people's actually employment conditions, is likely to have an even greater effect than other kinds of AI we've seen thus far. It's because, and we've seen this, we've seen companies that are digitally or AI native and have basically operated and get IPOs and the like with almost no employees. That is a hugely disruptive thing if you are a company that has largely worked on manual individual workers for many tasks and now want to be disruptive for agentic AI.

In this paper, we say today's clinicians who are seeing this coming down the pipeline, AI that's introduced with the goal of serving clients or patients in the case of clinicians, that clinicians themselves, their activities are being monitored, surveilled, and the like. And that this is a moment where they might think about what stakes and what authority should they have in this conversation. I think engagement is really important. They often say, "If you're not sitting at the table, you're on the menu." Right?

So this idea that we want to have some form of engagement with the workforce and maybe it's not a veto, but a form of engagement rather than ending up in a situation where the headlines or there's labor unrest depending on the kind of space you're looking at. And then I think we want some conception of quality assurance. So if you are an employer, a business introducing AI to manage your workforce or to surveil your workforce, I think it's reasonable for the workforce to say, "How do you know this is any good? If it's assisting you with employment decisions, what have you shown me to say that it doesn't have certain kinds of biases, that actually its outputs are verifiable, that you've kicked the tires on these sorts of things?"

So those are the things that I think of as the indisputable minimum. That is, a right to know you're dealing with AI, some forms of rights of participation to be involved at the table, and some kinds of assurances about the quality of the AI that's being used.

MCDONAGH-SMITH: Very, very, very interesting. Almost considerations of a social contract in a charter of the commons, but perhaps for the way in which we're deploying AI and how it affects each of the stakeholders.

Monday morning priorities

MCDONAGH-SMITH: So, Glenn, a final question for the senior business leaders here with us today on the program who are walking into their organizations on Monday morning. What are the one or two things you'd suggest that they make a priority? If you had their ear for, let's say, 90 seconds in an elevator, what would you tell them?

COHEN: Great question. Okay. So, two things I'll tell them. One on the front end, one on the back end. Okay. On the back end, I think it's this idea of post-implementation review. This idea that this is not a "set it and forget it moment," that if you're going to embark on any major integration of AI into what you do, you should have an analysis plan set up upfront and when financed upfront as to how you're going to evaluate how the AI is doing, how you're going to look for changes over time, how are you going to improve it, and also if necessary, how you're going to decommission it and to make sure you're willing to do that. Because again, if you're not willing to do it, to me, you can get yourself into a heap of trouble by just kind of implementing and hoping about it.

Okay. So that's on the back end, the planning for the back end. On the front end, think it's just to be very candid about what your goals are. That often, I think AI is presented as a shiny new toy and people come in, and you're getting pitched constantly both internally and externally about not being left behind and why you should have AI, in particular agentic AI integration.

And I think it's really just to look very, very carefully about what your goals are and ask yourself, you know, quite directly, "Is this something we think we can implement successfully or not?" And I do think that nexus between people can show you all sorts of incredibly impressive data regarding how it performs on verification data sets and new data and even your own data. But that's very different than knowing how it's going to play within your organization. So thinking about the realities of organizational dynamics and making sure that this is actually a solve for a problem you have rather than creating a new problem you don't have, I think is an important thing to think about.

Metrics and values

MCDONAGH-SMITH: If and when we're introducing new approaches, perhaps it makes sense to kind of consider, well, perhaps the legacy units of measurement and metrics of value that I'm inheriting here might not actually be capturing the impact and the value. Is there any guidance you would give to our business leaders around how they might fine-tune their perception so they're able to recognize where the impact and value is being created and, if necessary, identify new units of measurement?

COHEN: I love this question because often we say we want to measure what we value, but we end up valuing what we can measure. Right? So we're following the measurement

capacities and saying that's what we value instead of thinking about what it is we actually want to value. So here too, AI can be actually a major help in quantifying lots of data that you might have.

And again, one of the real revolutions of machine learning was to go from data that had to be extremely tightly structured for analysis, now have much more capacity for multimodal analysis and putting lots of data together. But that said, there has to be a values conversation at the heart of it. Right? It's to say, if you are a hospital system, right, do you care about catching the 1% of cancers you're missing? Do you care about expanding your ability to reach people in a rural population? Do you care about making more time for your physicians to avoid burnout? Those are all three very valuable things to do, but they require very different kind of strategies and also different kind of AI could be involved in all three, but the AI you would be designing would be completely different. But you can't really know if you have a problem until you've looked to see what the problem is.

And indeed, within an organization, there is often, I think, a kind of stoicism is to say, "We're fine. We can go forward. We'll be okay." Right? That if you talk to individual people. But having this more 360 view that sometimes can be the result of AI analysis can be extremely helpful in trying to get other kinds of metrics, not to replace those metrics in the human view about what the problem is, but to supplement it. My favorite example here is pharmacovigilance, right, is that we can have individual people report in adverse events for drugs, but we also can learn a lot from having AI essentially do some analysis of the information we already have.

MCDONAGH-SMITH: So, Glenn, thank you for the full range of insights and thoughts you've shared with us today. Fascinating and very useful. A big thank you from us all.

COHEN: Thank you so much for having me.

MCDONAGH-SMITH: Thank you.

SPEAKER: Did you understand all the concepts covered in this video? If you'd like to review any of the content, please click on the chapter menu.