




MIT SLOAN SCHOOL OF MANAGEMENT
MIT SCHWARZMAN COLLEGE OF COMPUTING



MODULE 2 UNIT 2
Media Set Video 2 Transcript

Module 2 Unit 2 Media Set Video 2 Transcript

Human-in-the-loop

PAUL MCDONAGH-SMITH: And so, Kanch, as agentic AI systems become more autonomous, what are the best practices that you see are emerging for managing risk, for managing bias, explainability, especially in those regulated or highly regulated industries?

KANCHANA PATLOLLA: So for me, regulated industries and autonomous do not go hand in hand at this point of time. Not saying it won't be there. Maybe even six months from now if you were to watch this recording, it might become outdated. I don't know. But, at this point of time, autonomous and regulated industries doesn't go hand in hand.

But when we are developing a product, we definitely need to like I specify this in all my conversations human-in-the-loop is still a must. The autonomous nature of it could be provided with some threshold. Like, hey, I don't care if it's like five dollars fraud and I'm going to let it go. That's something which we have always done. Even pre generative AI, we have created thresholds and we have handled it, I think we would continue to do it. But in general, your AI application needs to have a human-in-the-loop, at least at this point of time. Even if it's nonregulated, for me, that's important.

Then the next thing when you think about with respect to managing risk or bias or explainability, one of the things which has been a problem, even in our MLOps world where we were considering the evolution of CI-CD-CT-CM. We used to call continuous training, continuous monitoring, like that world. To now, where it's continuous evaluation. Evaluation should not stop in your just development phase because most organizations, you think of like, you know, comparing against multiple models, making sure it's getting to an accuracy they're looking for, doing some prompt-tuning or, you know, other providing more context and they're able to get to it. And there is evaluation with respect to benchmarks they are doing. They are having some golden truth evaluating. So these standard practices we have always known is still adopted in that regard.

But, like, in my opinion, the continuous evaluation of your results and the outcome for skews, needs to be inbuilt in your architecture. So it's less about actually agentic application. It's more about AgentOps. That's the new term which is coined, which is more about how do you build this deployment journey, which has the evaluation monitoring inbuilt into it, which identifies the skews of the results which comes in and notifies us. Because we know this, like, right, having built something with one model or version of a model and you upgrade it to a different version and everything is gone; it doesn't work. But in a production system, it's more seen. It's more obvious. So when I think about all of this topics, risk, bias, or explainability, the first thing is understanding your deployment process is the most important one for me.

SPEAKER: Did you understand all the concepts covered in this video? If you'd like to review any of the content, please click on the chapter menu.